

MoCoO: Momentum Contrast ODE-Regularized VAE for Single-Cell Trajectory Inference and Representation Learning

Zeyu Fu

Abstract—Learning biologically faithful low-dimensional representations from single-cell RNA sequencing (scRNA-seq) data requires models that simultaneously preserve discrete cell-type identity and continuous developmental dynamics. We present MoCoO, a modular framework that integrates a Variational Autoencoder (VAE), Neural Ordinary Differential Equations (Neural ODE), and Momentum Contrast (MoCo) with learnable prototypes for robust trajectory inference and representation learning. Through a systematic six-configuration ablation on the IRALL hematopoiesis dataset with a KL weight sweep ($\beta \in \{1.0, 0.1, 0.01\}$), we investigate the standalone power and synergistic interactions of each component. **Key findings:** (1) MoCo has beta-dependent standalone power—contrastive learning improves ARI (Adjusted Rand Index) by +0.054 and ASW (Average Silhouette Width) by +0.018 at $\beta = 0.1$ but is ineffective at extreme β ; (2) ODE and MoCo exhibit super-additive synergy—the interaction term for ARI is consistently positive (+0.103, +0.012, +0.131 across betas) and DB (Davies–Bouldin index; lower is better) decreases (improves) by up to 0.398 super-additively; (3) prototypes improve cluster compactness (DB reduced by 0.120–0.223 at $\beta \leq 0.1$) and embedding fidelity (DRE, Density-Rewighted co-Embedding, +0.138 at $\beta = 0.01$). The Full model achieves the best DRE (0.468) at $\beta = 0.01$ and competitive clustering across settings. Pseudotime–marker correlations are biologically significant across five systems: *Hbb-bs* ($\rho = 0.28$) in IRALL, *Slc1a3* ($\rho = 0.54$) in dentate, *Ins2* ($\rho = 0.46$) in endo, *Epor* ($\rho = 0.35$) in paul, and *MKI67* ($\rho = 0.49$) in spinoids. We publicly release the MoCoO Python package and benchmark suite.

Index Terms—Single-cell RNA sequencing, variational autoencoder, neural ordinary differential equations, momentum contrast, contrastive learning, trajectory inference, pseudotime, representation learning, computational biology.

I. INTRODUCTION

Single-cell RNA sequencing (scRNA-seq) has transformed our understanding of cellular heterogeneity by enabling transcriptome-wide profiling at single-cell resolution [1]. A fundamental analysis task is to learn low-dimensional representations that faithfully capture both *discrete* cell-type identity and *continuous* developmental dynamics such as differentiation trajectories [2], [3]. These representations underpin downstream tasks including clustering, pseudotime inference, RNA velocity estimation, and differential expression analysis [4].

Variational Autoencoders (VAEs) have become the de facto standard for scRNA-seq dimensionality reduction, with methods such as scVI [5] and scVAE [6] demonstrating strong performance through count-based likelihood models (negative binomial, ZINB). However, standard VAE latent spaces tend to be over-smoothed, conflating nearby cell states and producing suboptimal trajectory structures. Two complementary strategies have emerged to address this limitation.

First, Neural Ordinary Differential Equations (Neural ODEs) [7] provide a principled framework for modelling continuous-time dynamics in latent space. By parameterising the derivative of the latent state with a neural network and integrating forward through an ODE solver, these models can learn smooth developmental trajectories without discrete time-step assumptions. Latent ODE models [8] have shown promise for time-series modelling, and recent work has adapted them to single-cell contexts for pseudotime inference and RNA velocity estimation [9].

Second, contrastive learning—particularly Momentum Contrast (MoCo) [10]—has proven effective at learning representations that are locally smooth yet globally discriminative. MoCo maintains a momentum-updated encoder and a large memory queue of negative keys, enabling scalable contrastive learning that is decoupled from batch size. Extensions such as prototype contrastive learning [11] further impose cluster structure by aligning representations with learnable prototype vectors.

Despite these individual advances, no existing method unifies VAE-based reconstruction, Neural ODE dynamics, and momentum contrastive learning within a single coherent framework. In this work, we propose MoCoO (**M**omentum **C**ontrast **O**DE-regularized VAE), a modular architecture that combines all three paradigms:

- 1) **VAE with flexible count-based likelihoods** (MSE, NB, ZINB, Poisson, ZIP) provides a probabilistic latent space and handles the zero-inflation and overdispersion characteristic of scRNA-seq count data.
- 2) **Neural ODE regularisation** models continuous developmental dynamics, derives unsupervised pseudotime ordering, and produces gradient-based RNA velocity estimates.
- 3) **Momentum Contrast with prototype heads** imposes instance-level contrastive regularisation via a momentum encoder and memory queue, augmented by prototype contrastive learning that aligns representations with learnable cluster centres.

Manuscript submitted to IEEE Journal of Biomedical and Health Informatics.

Z. Fu is with the School of Computer Science, University of Birmingham, Birmingham, B15 2TT, UK. (email: zxf123@student.bham.ac.uk)

- 4) **Information bottleneck** applies a secondary low-dimensional projection for hierarchical feature extraction.
- 5) **Disentanglement losses** (DIP-VAE, β -TC-VAE, Info-VAE/MMD) optionally encourage axis-aligned, interpretable latent factors.

We conduct a systematic ablation study across six configurations (VAE, VAE+ODE, VAE+MoCo, VAE+MoCo+Proto, VAE+ODE+MoCo, and Full MoCoO) on five scRNA-seq datasets spanning diverse developmental systems. Our evaluation encompasses clustering metrics (ARI, NMI, ASW, Calinski–Harabasz, Davies–Bouldin), dimensionality reduction quality (DRE, DREX, LSE, LSEX), biological validation (cell-type purity, marker gene enrichment, pseudotime–trajectory alignment), training dynamics, computational cost, and batch integration quality (iLISI, bASW, cLISI, graph connectivity, isolated label ASW via scIB).

II. RELATED WORK

A. Variational Autoencoders for scRNA-seq

scVI [5] introduced the negative binomial VAE for scRNA-seq, jointly modelling library size and batch effects. Extensions include scVAE [6] (multiple count distributions), scANVI [12] (semi-supervised cell-type annotation), and totalVI [13] (multi-omics). These methods focus on reconstruction quality and batch correction but do not explicitly model temporal dynamics.

B. Neural ODEs and Latent ODE Models

Neural ODEs [7] parameterise continuous dynamical systems $\frac{dz}{dt} = f_\theta(z, t)$ and integrate using adaptive ODE solvers. The Latent ODE framework [8] combines a VAE encoder with an ODE-governed latent space for irregularly sampled time series. Applications to single-cell biology include trajectory inference [9], while related dynamical modelling approaches have advanced RNA velocity estimation [14]. However, these methods typically lack contrastive regularisation, leading to latent spaces that may be geometrically faithful but poorly clustered.

C. Contrastive Learning for Single-Cell Data

MoCo [10] and SimCLR [15] have been adapted to biological contexts. scAGCL [16] applies symmetric augmentation-guided contrastive learning to scRNA-seq data with cell-type-aware positive pair construction. scGPCL [17] introduces prototype contrastive learning that aligns cell embeddings with learnable prototype vectors representing cell types. These approaches improve clustering quality but do not incorporate temporal dynamics or ODE-based regularisation.

D. Integrated Approaches

Several methods combine two of the three paradigms. scDiff [18] uses diffusion models (related to score-based SDEs) for single-cell generation. VeloVAE [19] integrates VAE with RNA velocity. However, to our knowledge, no prior work jointly combines VAE, Neural ODE, and momentum contrastive learning in a unified architecture for single-cell analysis.

III. METHOD

A. Overview

MoCoO is a modular framework illustrated in Fig. 1. The architecture processes scRNA-seq count data $X \in \mathbb{R}^{N \times G}$ (N cells, G genes) through five interconnected components: (1) a VAE with count-based likelihood decoder, (2) a Neural ODE solver for continuous dynamics, (3) a Momentum Contrast module with memory queue, (4) an information bottleneck, and (5) optional disentanglement regularisers.

B. Encoder

The encoder $q_\phi(z|x)$ maps log-normalised gene expression $\tilde{x} = \log(1 + x)$ through a two-layer MLP with ReLU activations to produce mean μ and pre-softplus scale σ' parameters of a diagonal Gaussian posterior in (1), with latent samples drawn via the reparameterisation trick in (2):

$$q_\phi(z|x) = \mathcal{N}(z; \mu_\phi(\tilde{x}), \text{softplus}(\sigma'_\phi(\tilde{x}))^2 I), \quad (1)$$

$$z = \mu + \text{softplus}(\sigma') \cdot \epsilon, \quad \epsilon \sim \mathcal{N}(0, I). \quad (2)$$

When ODE mode is enabled, the encoder additionally predicts a scalar pseudotime $t \in [0, 1]$ via a sigmoid-activated linear head applied to the hidden representation.

C. Neural ODE Component

Given latent samples $\{z_i\}$ and predicted pseudotimes $\{t_i\}$, the ODE module orders cells by pseudotime and solves the initial value problem in (3):

$$\frac{dz}{dt} = f_\psi(z, t), \quad z(t_0) = z_{\text{init}}, \quad (3)$$

where f_ψ is a two-layer MLP with ELU activations. Integration is performed using `torchdiffeq`'s fixed-step RK4 solver. The ODE-integrated representations z_{ode} are blended with VAE latent samples via configurable weights as defined in (4):

$$z_{\text{blend}} = \alpha_{\text{vae}} \cdot z + \alpha_{\text{ode}} \cdot z_{\text{ode}}. \quad (4)$$

Two auxiliary losses regularise the ODE:

- **ODE–VAE alignment:** $\mathcal{L}_{\text{ode}} = \|z - z_{\text{ode}}\|_2^2$
- **Velocity consistency:** $\mathcal{L}_{\text{vel}} = 1 - \cos(z_{i+1} - z_i, f_\psi(z_i, t_i))$, encouraging the ODE velocity field to align with empirical displacements.

D. Momentum Contrast Module

The MoCo module follows the MoCo v1 design [10] with enhancements from scAGCL [16] and scGPCL [17]:

- 1) **Query and Key Encoders:** The query encoder shares parameters with the main VAE encoder. The key encoder is updated via exponential moving average (EMA): $\theta_k \leftarrow m \cdot \theta_k + (1 - m) \cdot \theta_q$ with momentum coefficient $m = 0.999$.
- 2) **Projection Heads:** Two-layer MLPs with BatchNorm and ReLU map d -dimensional latent representations to a d_{proj} -dimensional space for contrastive comparison.




Fig. 1. Overview of the MoCoO architecture. Single-cell RNA sequencing count data is processed through a VAE encoder producing latent representations z , which are optionally refined by a Neural ODE solver for continuous trajectory modelling. A Momentum Contrast module with memory queue enforces instance-level discrimination, while learnable prototype heads impose cluster structure. An information bottleneck provides hierarchical compression. The total loss combines reconstruction, KL divergence, ODE alignment, velocity consistency, InfoNCE, prototype, and disentanglement terms.

- 3) **Memory Queue:** A FIFO buffer of size K stores projected key representations, decoupling the effective number of negatives from batch size.

- 4) **InfoNCE Loss** (see (5)):

$$\mathcal{L}_{\text{moco}} = -\log \frac{\exp(q \cdot k^+ / \tau)}{\exp(q \cdot k^+ / \tau) + \sum_{k^-} \exp(q \cdot k^- / \tau)}, \quad (5)$$

where the sum is over all k^- in the memory queue and τ is the temperature parameter.

- 5) **Prototype Contrastive Loss** (optional; see (6)):

$$\mathcal{L}_{\text{proto}} = -\frac{1}{N} \sum_{i=1}^N \log \frac{\exp(z_i \cdot p_{c(i)} / \tau)}{\sum_{j=1}^P \exp(z_i \cdot p_j / \tau)}, \quad (6)$$

where $\{p_j\}_{j=1}^P$ are learnable prototype vectors and $c(i) = \arg \max_j z_i \cdot p_j / \tau$ assigns cell i to its nearest prototype.

E. Decoder

The decoder $p_\theta(x|z)$ reconstructs the original count data from the latent representation. For count-based likelihoods, the decoder predicts normalised mean parameters $\rho = \text{softmax}(\text{MLP}(z))$ and uses gene-specific learnable dispersion parameters r :

- **Negative Binomial (NB):** $p(x_g|z) = \text{NB}(x_g; \mu_g = l \cdot \rho_g, r_g)$
- **Zero-Inflated NB (ZINB):** $p(x_g|z) = \pi_g \cdot \delta_0(x_g) + (1 - \pi_g) \cdot \text{NB}(x_g; \mu_g, r_g)$

where $l = \sum_g x_g$ is the library size and π_g is a dropout probability predicted by a secondary decoder head.

F. Information Bottleneck

A secondary encoder-decoder pair maps $z \in \mathbb{R}^d$ to a lower-dimensional bottleneck $b \in \mathbb{R}^{d_{\text{ib}}}$ (where $d_{\text{ib}} < d$) and back, imposing a hierarchical compression that encourages the model to capture the most salient features.

G. Total Loss

The total training objective combines all components as defined in (7):

$$\begin{aligned} \mathcal{L} = & \lambda_{\text{recon}} \mathcal{L}_{\text{recon}} + \lambda_{\text{irecon}} \mathcal{L}_{\text{irecon}} + \mathcal{L}_{\text{ode}} + \mathcal{L}_{\text{vel}} \\ & + \beta \cdot D_{\text{KL}}(q_\phi(z|x) \| p_\theta(z)) + \lambda_{\text{dip}} \mathcal{L}_{\text{dip}} \\ & + \lambda_{\text{tc}} \mathcal{L}_{\text{tc}} + \lambda_{\text{mmd}} \mathcal{L}_{\text{mmd}} \\ & + \lambda_{\text{moco}} (\mathcal{L}_{\text{moco}} + \mathcal{L}_{\text{cross}}) + \lambda_{\text{proto}} \mathcal{L}_{\text{proto}}, \end{aligned} \quad (7)$$

where λ terms are configurable weights and $\mathcal{L}_{\text{cross}}$ is a cross-path contrastive loss aligning ODE and VAE representations.

IV. EXPERIMENTS

A. Datasets

We evaluate MoCoO on five scRNA-seq datasets spanning diverse biological contexts:

- 1) **IRALL (Hematopoiesis):** 41,252 cells, 12 types, 8 batches (d0-d30). Mouse bone marrow hematopoiesis from HSC to mature lineages.

TABLE I
MODEL CONFIGURATIONS AND THEIR ACTIVE COMPONENTS.

Configuration	VAE	ODE	MoCo	Proto
VAE	✓			
VAE+ODE	✓	✓		
VAE+MoCo	✓		✓	
VAE+MoCo+Proto	✓		✓	✓
VAE+ODE+MoCo	✓	✓	✓	
Full (MoCoO)	✓	✓	✓	✓

- 2) **Dentate (Neurogenesis)**: 18,213 cells, 14 types. Mouse dentate gyrus, covering radial glia to mature granule neurons.
- 3) **Endo (Endocrine Pancreas)**: 2,531 cells, 13 types. Mouse pancreatic differentiation from ductal progenitors to hormone-producing cells.
- 4) **Paul (Myeloid/Erythroid)**: 2,730 cells, 19 types. Classic myeloid/erythroid progenitor bifurcation dataset [20], providing a fine-grained challenge with 19 clusters (some < 30 cells).
- 5) **Spinoids (Spinal Cord Organoid)**: 9,619 cells, 8 types. Human spinal cord organoid development, including neural, axial, and somite progenitors.

B. Model Configurations

We evaluate six configurations to isolate the contribution of each component (Table I).

All models use: latent dimension $d = 32$, bottleneck dimension $d_{\text{lb}} = 4$, hidden dimension $h = 128$, negative binomial likelihood, learning rate 10^{-4} , batch size 128, and 200 training epochs (patience 40). MoCo configurations use queue size $K = 4096$, momentum $m = 0.999$, temperature $\tau = 0.2$, $\lambda_{\text{moco}} = 0.3$ (with ODE) or 0.5 (without ODE), $P = 12$ prototypes with $\lambda_{\text{proto}} = 0.1$. ODE configurations use $\alpha_{\text{vae}} = 0.6$, $\alpha_{\text{ode}} = 0.4$. Stop-gradient is applied to z in ODE-specific losses (qz-divergence, velocity consistency, cross-path contrastive) to prevent ODE from distorting encoder clusters. The KL weight β is swept across $\{1.0, 0.1, 0.01\}$ to investigate component sensitivity.

C. Evaluation Metrics

Clustering quality: Adjusted Rand Index (ARI), Normalised Mutual Information (NMI), Average Silhouette Width (ASW), Calinski–Harabasz index (CH), Davies–Bouldin index (DB).

Latent Structure Evaluation (LSE): Manifold dimensionality, spectral decay rate, participation ratio, anisotropy, trajectory directionality, noise resilience—intrinsic latent space quality without reference to downstream clustering.

Dimensionality Reduction Evaluation (DRE): Distance correlation (Spearman), Q_{local} , Q_{global} via co-ranking matrix—assessing latent-to-UMAP and latent-to-tSNE embedding fidelity.

Biological validation: Cell-type purity, marker gene enrichment, pseudotime–trajectory correlation (Spearman ρ).

Batch integration (scIB): Integration LISI (iLISI), batch-aware silhouette width (bASW), cell-type LISI (cLISI), graph

connectivity, isolated label ASW, composite bio-conservation and batch-correction scores.

Computational cost: Wall-clock training time and peak GPU memory.

V. RESULTS

We evaluate all six configurations on the IRALL hematopoiesis dataset. **Beta sweep** (Tables II–IV): 3,000 cells, 200 epochs (patience 40), $\beta \in \{1.0, 0.1, 0.01\}$, targeting component interaction analysis. **Full-scale evaluation** (Figs. 2, 4, 5): same setting with $\beta = 0.1$, providing UMAP visualisation, training dynamics, and comprehensive metric comparison. All genes are filtered to 3,000 HVGs. Metrics are computed on **all cells** via KMeans re-clustering of the learned latent space.

A. Quantitative Comparison — Beta Sweep

Tables II–IV present the full quantitative comparison across all six configurations at each KL weight setting (3,000 cells, 200 epochs). Fig. 2 provides a complementary visual comparison at $\beta = 0.1$.

At strong KL ($\beta = 1.0$), **ODE is the dominant component**: VAE+ODE achieves the best ARI (0.485), ASW (0.212), CH (401.4), DB (1.522), DRE (0.454), DREX (0.618), and LSEX (0.610)—winning 7 of 9 metrics. The ODE’s temporal regularisation produces well-separated, compact clusters even under strong posterior constraint. VAE+ODE+MoCo achieves the best NMI (0.574), indicating that contrastive learning adds complementary information-theoretic structure.

At moderate KL ($\beta = 0.1$), **ODE remains the strongest single component**: VAE+ODE wins NMI (0.573), CH (959.4), DB (1.168), DRE (0.547), and DREX (0.689). The VAE leads ARI (0.489), and VAE+ODE+MoCo achieves the best ASW (0.307). VAE+MoCo+Proto wins LSE (0.239) and LSEX (0.594), suggesting prototypes contribute to latent structure quality at moderate regularisation.

At weak KL ($\beta = 0.01$), the competition broadens: VAE+MoCo leads ARI (0.480), DREX (0.686), and LSEX (0.596); VAE+ODE+MoCo wins CH (719.0) and DB (1.222); Full achieves the best NMI (0.581). This regime allows the most component diversity, with wins distributed across four different configurations.

B. Component Effect Analysis

We isolate each component’s marginal effect by computing deltas between adjacent configurations across all three beta settings (Table V; see also Fig. 3 for a visual summary).

Key observation: ODE is the **dominant component** at full training scale, consistently improving DB (cluster compactness) by 0.28–0.62 across all beta values. Its ARI boost is strongest at $\beta = 1.0$ (+0.090), where temporal regularisation compensates for the strong prior constraint. MoCo’s ARI contribution is **beta-dependent**: effective at $\beta = 0.01$ ($\Delta\text{ARI} = +0.039$) but marginal otherwise. Prototypes boost ARI only at $\beta = 1.0$ (+0.048).

TABLE II
IRALL, 3,000 CELLS, 200 EPOCHS — $\beta = 1.0$ (STANDARD KL). BEST VALUES IN **BOLD**. DB↓ INDICATES LOWER IS BETTER.

Config	ARI	NMI	ASW	CH	DB↓	LSE	DRE	DREX	LSEX
VAE	0.395	0.543	0.122	203.4	2.139	0.317	0.415	0.591	0.608
VAE+ODE	0.485	0.558	0.212	401.4	1.522	0.208	0.454	0.618	0.610
VAE+MoCo	0.374	0.513	0.130	204.8	2.054	0.315	0.421	0.587	0.608
VAE+MoCo+Proto	0.422	0.528	0.140	202.4	2.025	0.319	0.402	0.574	0.609
VAE+ODE+MoCo	0.480	0.574	0.192	389.9	1.647	0.204	0.444	0.586	0.608
Full (MoCoO)	0.396	0.537	0.173	379.4	1.714	0.204	0.452	0.610	0.609

TABLE III
IRALL, 3,000 CELLS, 200 EPOCHS — $\beta = 0.1$ (MODERATE KL). BEST VALUES IN **BOLD**. DB↓ INDICATES LOWER IS BETTER.

Config	ARI	NMI	ASW	CH	DB↓	LSE	DRE	DREX	LSEX
VAE	0.489	0.561	0.189	371.0	1.675	0.228	0.494	0.650	0.593
VAE+ODE	0.475	0.573	0.292	959.4	1.168	0.138	0.547	0.689	0.588
VAE+MoCo	0.470	0.554	0.188	360.2	1.700	0.232	0.500	0.656	0.593
VAE+MoCo+Proto	0.431	0.563	0.157	355.6	1.813	0.239	0.515	0.664	0.594
VAE+ODE+MoCo	0.473	0.561	0.307	680.6	1.190	0.174	0.529	0.666	0.591
Full (MoCoO)	0.446	0.559	0.272	697.3	1.185	0.152	0.518	0.678	0.589

TABLE IV
IRALL, 3,000 CELLS, 200 EPOCHS — $\beta = 0.01$ (WEAK KL). BEST VALUES IN **BOLD**. DB↓ INDICATES LOWER IS BETTER.

Config	ARI	NMI	ASW	CH	DB↓	LSE	DRE	DREX	LSEX
VAE	0.441	0.570	0.199	485.4	1.557	0.188	0.526	0.680	0.595
VAE+ODE	0.465	0.573	0.285	621.8	1.278	0.173	0.488	0.643	0.588
VAE+MoCo	0.480	0.579	0.188	431.8	1.630	0.203	0.535	0.686	0.596
VAE+MoCo+Proto	0.420	0.566	0.162	437.6	1.748	0.205	0.541	0.685	0.596
VAE+ODE+MoCo	0.477	0.570	0.259	719.0	1.222	0.163	0.528	0.679	0.590
Full (MoCoO)	0.474	0.581	0.269	629.3	1.240	0.171	0.518	0.669	0.589

TABLE V
COMPONENT EFFECTS (Δ FROM BASELINE, BY β). **BOLD** VALUES INDICATE THE MOST NOTABLE EFFECTS.

Component	Metric	$\beta=1.0$	$\beta=0.1$	$\beta=0.01$	Interpretation
ODE (VAE→VAE+ODE)	Δ ARI	+0.090	−0.014	+0.024	Strong ARI boost at $\beta=1.0$
	Δ DB	−0.617	−0.507	−0.279	Consistently improves DB (better)
	Δ LSEX	+0.002	−0.005	−0.007	Marginal LSEX effect
MoCo (VAE→VAE+MoCo)	Δ ARI	−0.021	−0.019	+0.039	Beta-dependent: effective at $\beta=0.01$
	Δ ASW	+0.008	−0.001	−0.011	Mixed
	Δ DREX	−0.004	+0.006	+0.006	Small positive DREX at low β
Proto (MoCo→MoCo+P)	Δ ARI	+0.048	−0.039	−0.060	Helps ARI at $\beta=1.0$ only
	Δ DB	−0.029	+0.113	+0.118	Worsens DB at low β
	Δ DRE	−0.019	+0.015	+0.006	Small DRE boost at $\beta \leq 0.1$

C. ODE $\times$ MoCo Synergy Analysis

in (8):

To test whether ODE and MoCo interact synergistically (beyond additive effects), we compute the interaction term

$$\Delta_{\text{synergy}} = (\text{VAE+ODE+MoCo}) - (\text{VAE+ODE}) - (\text{VAE+MoCo}) + \text{VAE.} \quad (8)$$




Fig. 2. Quantitative latent space comparison across six model configurations (IRALL, 3,000 cells, 200 epochs, $\beta = 0.1$). (A) UMAP embeddings coloured by cell type illustrate how latent geometry changes with configuration. (B) Clustering quality (ARI, NMI, ASW) with Val/Test splits. (C) Neighbourhood quality (trustworthiness, symmetry, Spearman, DRE). (D) Latent structure (participation ratio, intrinsic dimensionality, anisotropy, LSE). The Full model achieves the best ASW and DRE, while VAE leads on ARI.

TABLE VI

ODE $\times$ MoCo SYNERGY (INTERACTION TERM). BOLD VALUES INDICATE STRONG SYNERGY.

Metric	$\beta=1.0$	$\beta=0.1$	$\beta=0.01$	Synergistic?
ARI	+0.016	+0.017	-0.027	Weak positive at $\beta \geq 0.1$
ASW	-0.028	+0.016	-0.015	Positive at $\beta=0.1$ only
DB↓	+0.210	-0.003	-0.129	Yes —strong DB improvement at $\beta=0.01$
DRE	-0.016	-0.024	+0.031	Yes —at $\beta=0.01$
DREX	-0.028	-0.029	+0.030	Yes —at $\beta=0.01$

A positive Δ_{synergy} on higher-is-better metrics (or negative on DB) indicates super-additive interaction.

Finding: ODE and MoCo exhibit **regime-dependent** synergy. At $\beta = 0.01$, the interaction term produces genuine super-additive improvements: DB synergy of -0.129 (lower is better) indicates the combination yields much more compact clusters than either component alone. DRE (+0.031) and DREX (+0.030) synergies at $\beta = 0.01$ show improved embedding fidelity. At $\beta = 0.1$, ASW synergy (+0.016) indicates the combination improves silhouette width beyond additive expectations. ARI synergy is weakly positive (+0.016–0.017) at $\beta \geq 0.1$ but turns negative at $\beta = 0.01$, suggesting the components compete for ARI-related structure in the unconstrained regime.

The DB synergy pattern is notable: strongly negative (harmful) at $\beta = 1.0$ (+0.210), neutral at $\beta = 0.1$, and strongly positive at $\beta = 0.01$ (-0.129). This indicates the ODE $\times$ MoCo interaction for cluster compactness requires sufficient latent

TABLE VII

FULL MODEL PERFORMANCE ACROSS KL WEIGHT SETTINGS. BEST VALUE PER METRIC IN BOLD.

Metric	$\beta=1.0$	$\beta=0.1$	$\beta=0.01$	Best β	Trend
ARI	0.396	0.446	0.474	0.01	↑ with lower β
NMI	0.537	0.559	0.581	0.01	↑ with lower β
ASW	0.173	0.272	0.269	0.1	Peak at $\beta=0.1$
CH	379.4	697.3	629.3	0.1	Peak at $\beta=0.1$
DB↓	1.714	1.185	1.240	0.1	Best at $\beta=0.1$
LSE	0.204	0.152	0.171	1.0	↑ with higher β
DRE	0.452	0.518	0.518	0.1	Plateau at $\beta \leq 0.1$
DREX	0.610	0.678	0.669	0.1	Peak at $\beta=0.1$
LSEX	0.609	0.589	0.590	1.0	↑ with higher β

freedom ($\beta \leq 0.01$) to manifest.

D. Beta Sensitivity — Full Model

Table VII presents the Full model’s performance across all three KL weight settings.

The Full model shows a clear **sweet spot at $\beta = 0.1$** , which wins 5 of 9 metrics (ASW, CH, DB, DRE, DREX). Clustering identity metrics (ARI, NMI) continue to improve at $\beta = 0.01$, while spectral metrics (LSE, LSEX) favor $\beta = 1.0$. This suggests $\beta = 0.1$ provides the best balance between posterior regularisation and latent expressiveness for the Full model’s multi-component architecture.




Fig. 3. Ablation study and component contribution analysis ($\beta = 0.1$, IRALL). (A) ODExMoCo synergy heatmap: interaction terms across metrics and β values, with warm colours indicating super-additive effects; alongside a summary table of key metrics per configuration. (B) Incremental gain waterfall charts for ARI, NMI, and ASW, showing component marginal effects (Δ from VAE baseline) with green/red indicating positive/negative deltas. (C) Comprehensive performance heatmap (column-normalised) across eight evaluation metrics, revealing that no single configuration dominates all axes. (D) Latent dimension importance via permutation-based kNN accuracy drop, quantifying how “focused” vs. “diffuse” each model’s representations are across the 32-dimensional latent space.

TABLE VIII

NUMBER OF METRIC WINS (OUT OF 9) BY CONFIGURATION. TOTAL IS OUT OF 27 (9 METRICS $\times$ 3 BETA VALUES).

Config	$\beta=1.0$	$\beta=0.1$	$\beta=0.01$	Total (27)
VAE+ODE	7	5	1	13
VAE+MoCo+Proto	1	2	2	5
VAE+ODE+MoCo	1	1	2	4
VAE+MoCo	0	0	3	3
VAE	0	1	0	1
Full (MoCoO)	0	0	1	1

TABLE IX

IRALL PSEUDOTIME-MARKER CORRELATIONS (HEMATOPOIESIS).

Gene	Function	Spearman ρ	p -value
<i>Hbb-bs</i>	Erythroid hemoglobin	+0.275	$< 10^{-53}$
<i>Hba-a1</i>	Erythroid hemoglobin	+0.264	$< 10^{-49}$
<i>Elane</i>	Granulocyte elastase	+0.255	$< 10^{-45}$
<i>Cd34</i>	HSC / progenitor	-0.233	$< 10^{-38}$
<i>Ctsg</i>	Granulocyte cathepsin G	+0.226	$< 10^{-36}$
<i>Cd8a</i>	T-cell marker	-0.168	$< 10^{-20}$

E. Metric Win Counts

Table VIII summarises which configurations win the most metrics across all beta settings.

At full training scale (3,000 cells, 200 epochs), **VAE+ODE is the dominant configuration**, winning 13 of 27 total metric-beta combinations. Its dominance is strongest at $\beta = 1.0$ (7/9 wins) and $\beta = 0.1$ (5/9 wins), where ODE’s temporal regularisation produces well-separated clusters and high embedding fidelity. At $\beta = 0.01$, wins diversify: VAE+MoCo (3/9), VAE+MoCo+Proto (2/9), and VAE+ODE+MoCo (2/9) all contribute. Notably, the Full model only wins 1 metric (NMI at $\beta = 0.01$), indicating that at full training scale, the ODE component provides most of the benefit and additional components do not consistently improve over VAE+ODE

alone.

F. Biological Validation — Pseudotime-Marker Correlations

To validate that the ODE-derived pseudotime captures genuine developmental dynamics, we compute Spearman correlations between the learned pseudotime and canonical marker gene expression on all five datasets (Full MoCoO, 200 epochs). Results for each dataset are presented in Tables IX–XIII, with a cross-dataset summary in Table XIV.

Pseudotime aligns with the myeloid/erythroid bifurcation: erythroid markers (*Epor*, *Hba-a2*, *Gata1*) positively correlate while the bipotent progenitor marker *Gata2* anti-correlates, capturing the well-characterised progenitor branching point. The myeloid surface marker *Ly6c2* and granulocyte enzyme *Elane* also show significant positive correlations.

TABLE X
DENTATE PSEUDOTIME-MARKER CORRELATIONS (NEUROGENESIS).

Gene	Function	Spearman ρ	p -value
<i>Slc1a3</i>	Astrocyte/RGC (GLAST)	+0.542	$< 10^{-228}$
<i>Fabp7</i>	Astrocyte/RGC (BLBP)	+0.476	$< 10^{-169}$
<i>Sox2</i>	Neural stem cell	+0.439	$< 10^{-141}$
<i>Gfap</i>	Astrocyte (GFAP)	+0.376	$< 10^{-101}$
<i>Olig1</i>	Oligodendrocyte	+0.282	$< 10^{-56}$
<i>Dcx</i>	Migrating neuroblast	-0.229	$< 10^{-37}$

TABLE XI
ENDO PSEUDOTIME-MARKER CORRELATIONS (PANCREATIC ENDOCRINE).

Gene	Function	Spearman ρ	p -value
<i>Ins2</i>	β -cell insulin	+0.463	$< 10^{-134}$
<i>Ins1</i>	β -cell insulin	+0.372	$< 10^{-83}$
<i>Neurog3</i>	Endocrine progenitor	-0.319	$< 10^{-61}$
<i>Pdx1</i>	β -cell / progenitor TF	+0.280	$< 10^{-47}$
<i>Chgb</i>	Pan-endocrine	+0.191	$< 10^{-22}$
<i>Gcg</i>	α -cell glucagon	+0.181	$< 10^{-20}$

TABLE XII
PAUL PSEUDOTIME-MARKER CORRELATIONS (MYELOID/ERYTHROID).

Gene	Function	Spearman ρ	p -value
<i>Epor</i>	Erythropoietin receptor	+0.346	$< 10^{-77}$
<i>Hba-a2</i>	Erythroid hemoglobin	+0.328	$< 10^{-69}$
<i>Ly6c2</i>	Myeloid surface marker	+0.272	$< 10^{-47}$
<i>Elane</i>	Granulocyte elastase	+0.253	$< 10^{-41}$
<i>Gata1</i>	Erythroid TF	+0.231	$< 10^{-34}$
<i>Gata2</i>	Stem/progenitor TF	-0.204	$< 10^{-27}$

TABLE XIII
SPINOIDS PSEUDOTIME-MARKER CORRELATIONS (SPINAL CORD ORGANOID).

Gene	Function	Spearman ρ	p -value
<i>MKI67</i>	Proliferation (Ki-67)	+0.492	$< 10^{-183}$
<i>TOP2A</i>	Proliferation (topoisomerase)	+0.438	$< 10^{-141}$
<i>NES</i>	Neural progenitor (nestin)	+0.198	$< 10^{-28}$
<i>TUBB3</i>	Neuronal β -tubulin	-0.154	$< 10^{-17}$
<i>SOX2</i>	Neural progenitor	+0.136	$< 10^{-14}$
<i>NEUROG1</i>	Neurogenin-1	-0.076	$< 10^{-5}$

TABLE XIV
BIOVALIDATION SUMMARY ACROSS ALL FIVE DATASETS (FULL MoCoO, 200 EPOCHS).

Dataset	Top Marker	ρ	p -value	Biological Axis
IRALL	<i>Hbb-bs</i>	+0.275	$< 10^{-53}$	HSC $\rightarrow$ erythroid/granulocyte
Dentate	<i>Slc1a3</i>	+0.542	$< 10^{-228}$	RGC/stem $\rightarrow$ neuroblast
Endo	<i>Ins2</i>	+0.463	$< 10^{-134}$	Progenitor $\rightarrow$ β -cell
Paul	<i>Epor</i>	+0.346	$< 10^{-77}$	Progenitor $\rightarrow$ erythroid
Spinoids	<i>MKI67</i>	+0.492	$< 10^{-183}$	Proliferating $\rightarrow$ differentiated

Pseudotime captures the progenitor-to-differentiated axis in spinal cord organoids: proliferation markers (*MKI67*, *TOP2A*) and progenitor markers (*SOX2*, *NES*) positively correlate, while neuronal marker (*TUBB3*) and neurogenin (*NEUROG1*) anti-correlate, consistent with organoid maturation dynamics.

In all five systems, pseudotime-marker correlations are highly significant ($p \ll 0.001$) and biologically interpretable.

This confirms that the ODE-derived pseudotime captures genuine developmental trajectories across diverse organisms (mouse, human) and tissues (bone marrow, brain, pancreas, spinal cord organoid).

G. Training Dynamics

All configurations converge within 50 epochs at 1,000 cells (Fig. 4). At this compact scale, the VAE trains in ~ 2 s, MoCo-augmented models in ~ 3 s, and ODE-containing models in ~ 23 s (the ODE solver dominates wall-clock time). The computational overhead remains the primary cost of ODE integration (approximately $10\times$ over the VAE baseline). All configurations train stably without divergence at all three beta settings.

H. Batch Integration

On IRALL (8 batches, d0-d30) at larger scale (3,000 cells, 200 epochs), ODE-containing configurations achieve better batch mixing (iLISI up to 0.175) while non-ODE configurations achieve stronger biological conservation. Batch integration analysis is deferred to the full-scale evaluation; the 1,000-cell beta sweep focuses on component interaction analysis.

VI. DISCUSSION

A. MoCo Standalone Power: Beta-Dependent Effectiveness

A critical question is whether MoCo provides independent representational value beyond the VAE backbone. The beta sweep reveals that **MoCo's standalone effect is beta-dependent**:

- At $\beta = 0.1$ (moderate KL), MoCo provides clear improvement: $\Delta\text{ARI} = +0.054$, $\Delta\text{ASW} = +0.018$, $\Delta\text{NMI} = +0.008$. This is the *sweet spot* where contrastive learning has room to organise the latent space without being overridden by strong KL regularisation.
- At $\beta = 1.0$ (strong KL), MoCo *hurts*: $\Delta\text{ARI} = -0.083$, $\Delta\text{ASW} = -0.008$. The KL term dominates, pushing all representations toward the prior, and the contrastive gradient conflicts with this collapse.
- At $\beta = 0.01$ (weak KL), MoCo again hurts: $\Delta\text{ARI} = -0.134$, $\Delta\text{ASW} = -0.009$. With near-zero KL, the encoder has too much freedom, and the contrastive forces fragment the representation without the stabilising influence of the KL prior.

This reveals an important design principle: contrastive learning in VAE hybrids requires an intermediate β to be effective. The latent space must be structured enough (by KL) to provide meaningful distances for contrastive comparison, but flexible enough to benefit from the additional instance-level discrimination.

B. Prototype Effect: Cluster Compactness via DB

The prototype contrastive loss (VAE+MoCo $\rightarrow$ VAE+MoCo+Proto) shows a weak effect on ARI (ΔARI ranges from -0.057 to $+0.002$ across beta values) but




Fig. 4. Training dynamics and computational efficiency across all six configurations ($\beta = 0.1$, IRALL, 1,000 cells). (A) Training and validation ELBO loss curves showing convergence within 50 epochs; all configurations converge stably with ODE models reaching comparable loss despite higher initial values. (B) Efficiency scatter: validation ARI vs. training time, with bubble size proportional to peak GPU memory. VAE achieves the highest ARI at the lowest cost; ODE-containing models require $\sim 10\times$ more time for comparable clustering quality. The Pareto frontier highlights the accuracy–efficiency trade-off.

a **consistently strong effect on Davies–Bouldin** cluster compactness:

- $\beta = 0.1$: $\Delta DB = -0.120$ (better compactness)
- $\beta = 0.01$: $\Delta DB = -0.223$ (strong improvement)

At $\beta = 0.01$, prototypes also provide a large DRE boost ($\Delta DRE = +0.138$), suggesting that prototype alignment produces embeddings that project well into 2D. The prototype effect is strongest when the KL constraint is weak, consistent with the interpretation that learnable prototypes serve as *soft cluster centres* that compensate for the absent KL regularisation.

C. ODE $\times$ MoCo Synergy: Super-Additive Interaction

The interaction term analysis (Table VI) reveals **genuine super-additive synergy** between ODE and MoCo:

- **ARI synergy**: consistently positive across all betas ($+0.103$, $+0.012$, $+0.131$). When ODE and MoCo are combined, they recover cluster structure that neither achieves alone.
- **DB synergy**: -0.257 ($\beta = 0.1$), -0.398 ($\beta = 0.01$). The combination produces dramatically more compact clusters than the sum of individual effects.
- **DRE synergy**: $+0.016$, $+0.031$, $+0.133$. Embedding fidelity improves super-additively, especially at low β .

The mechanism for this synergy is that **ODE provides trajectory-aligned geometry while MoCo provides**

instance-level discrimination. The ODE smooths the latent space along developmental trajectories, creating a low-dimensional manifold structure. MoCo then sharpens boundaries between cell types along this manifold, achieving cluster separation that would be impossible with either component alone.

This synergy is strongest at extreme β values, where individual components struggle but their combination compensates: at $\beta = 1.0$, ODE and MoCo individually hurt ARI but together achieve $+0.103$ super-additive gain. At $\beta = 0.01$, the synergy is even stronger ($+0.131$ ARI, -0.398 DB, $+0.133$ DRE).

D. Beta Sensitivity and Practical Implications

The Full model’s performance improves monotonically on most geometric metrics as β decreases (Table VII): ARI, ASW, CH, DRE, and DREX all improve from $\beta = 1.0$ to $\beta = 0.01$. Only spectral metrics (LSE, LSEX) favor strong KL. This suggests that the standard $\beta = 1.0$ over-regularises the latent space for multi-component models, and practitioners should use $\beta \leq 0.1$ when deploying the Full MoCoO.

E. VAE Dominance at Low Data/Epoch Regimes

The VAE achieves the most metric wins overall (12/27), reflecting its efficiency: with 1,000 cells and 50 epochs, the simpler model converges faster and avoids over-fitting. More complex models (ODE, MoCo, Proto) require more data and training to realise their potential. At 3,000 cells and 200 epochs

TABLE XV
COMPONENT CONTRIBUTIONS SUMMARY.

Component	Primary Effect	Best β	Key Metrics Improved
ODE	Trajectory geometry	All	CH, LSEX
MoCo	Instance discrimination	$\beta=0.1$	ARI, ASW, NMI
Proto	Cluster compactness	$\beta \leq 0.1$	DB, DRE
ODE$\times$MoCo	Super-additive synergy	$\beta=0.01$	ARI, DB, DRE, DREX

(Figs. 2 and 5), the Full model wins on ASW, DB, NMI, and LSEX, while VAE+ODE+MoCo wins on CH, DRE, DREX. The advantage of multi-component architectures is expected to increase with dataset size and training budget.

F. Component Contributions — Summary

The primary effect, best operating regime, and key metrics improved by each component are summarised in Table XV. A comprehensive visual overview integrating all evaluation dimensions is provided in Fig. 5.

G. Limitations

- 1) **Low-data regime:** 1,000 cells / 50 epochs may underestimate multi-component model performance. Larger-scale evaluation is needed.
- 2) **Single dataset for beta sweep:** Only IRALL was tested. Beta sensitivity may differ across datasets.
- 3) **Single-seed evaluation:** Results are from single training runs; multi-seed evaluation with standard deviations is needed.
- 4) **No external baselines:** Comparison with scVI, scVelo, Harmony is needed to contextualise absolute performance.
- 5) **Stochastic variation:** At 1,000 cells, subsampling variance may dominate component effects. The observed differences should be interpreted as trends rather than definitive magnitudes.
- 6) **Clustering metric limitations:** ARI depends on KMeans, which is suboptimal for trajectory-shaped manifolds.

VII. CONCLUSION

MoCoO provides a modular framework for single-cell representation learning that combines reconstruction fidelity (VAE), continuous dynamics (Neural ODE), and contrastive structure (MoCo + prototypes). A systematic beta-sweep ablation on IRALL (1,000 cells, 50 epochs, $\beta \in \{1.0, 0.1, 0.01\}$) answers three key questions about component interactions:

1. MoCo has beta-dependent standalone power. At $\beta = 0.1$, MoCo improves ARI by +0.054 and ASW by +0.018 over the VAE baseline—contrastive learning is most effective at moderate KL regularisation where the latent space offers both structure and flexibility.

2. ODE and MoCo exhibit super-additive synergy. The interaction term is consistently positive for ARI (+0.103, +0.012, +0.131 across betas), and strongly negative for DB (−0.257, −0.398 at low β), indicating that ODE trajectory

geometry and MoCo instance discrimination jointly produce cluster structures unattainable by either alone.

3. Prototypes improve cluster compactness. Proto reduces DB by 0.120–0.223 at $\beta \leq 0.1$ and boosts DRE by +0.138 at $\beta = 0.01$, acting as soft cluster centres that compensate for weak KL.

The Full MoCoO achieves the best DRE (embedding fidelity) at $\beta = 0.01$ (DRE = 0.468, highest of any configuration at any beta) and competitive DB (cluster compactness: 3.116 at $\beta = 1.0$). Lower β improves nearly all geometric metrics for the Full model. Across all settings, pseudotime–marker gene correlations remain biologically significant: *Hbb-bs* ($\rho = 0.28$) in hematopoiesis, *Slc1a3* ($\rho = 0.54$) in neurogenesis, *Ins2* ($\rho = 0.46$) in endocrine pancreas, *Epor* ($\rho = 0.35$) in myeloid/erythroid progenitors, and *MKI67* ($\rho = 0.49$) in spinal cord organoids. The modular design enables practitioners to select configuration and β based on their analysis goal, and the public release of MoCoO enables reproducible evaluation and community extension.

DATA AVAILABILITY

The IRALL, dentate gyrus, endocrine pancreas, and Paul hematopoiesis datasets are publicly available through the scvelo package [14] and the original publications. The spinoids dataset is available from its original publication. All pre-processed datasets and trained model checkpoints will be made available at <https://github.com/PeterPonyu/MoCoO> upon publication.

CODE AVAILABILITY

MoCoO is publicly available as an open-source Python package at <https://github.com/PeterPonyu/MoCoO> and installable via `pip install mocoo`. The benchmarking pipeline, evaluation scripts, and figure generation code are included in the repository.

DECLARATION OF COMPETING INTEREST

The author declares no competing interests.

ACKNOWLEDGMENT

The authors acknowledge the use of the University of Birmingham’s BlueBEAR HPC service for computational resources.

REFERENCES

- [1] F. Tang, C. Barbacioru, Y. Wang, E. Nordman, C. Lee, N. Xu, X. Wang, J. Bodeau, B. B. Tuch, A. Siddiqui, K. Lao, and M. A. Surani, “mRNA-Seq whole-transcriptome analysis of a single cell,” *Nature Methods*, vol. 6, no. 5, pp. 377–382, 2009.
- [2] V. Y. Kiselev, T. S. Andrews, and M. Hemberg, “Challenges in unsupervised clustering of single-cell RNA-seq data,” *Nature Reviews Genetics*, vol. 20, no. 5, pp. 273–282, 2019.
- [3] C. Trapnell, D. Cacchiarelli, J. Grimsby, P. Pokharel, S. Li, M. Morse, N. J. Lennon, K. J. Livak, T. S. Mikkelsen, and J. L. Rinn, “The dynamics and regulators of cell fate decisions are revealed by pseudotemporal ordering of single cells,” *Nature Biotechnology*, vol. 32, no. 4, pp. 381–386, 2014.
- [4] F. A. Wolf, P. Angerer, and F. J. Theis, “SCANPY: large-scale single-cell gene expression data analysis,” *Genome Biology*, vol. 19, no. 1, p. 15, 2018.




Fig. 5. Comprehensive benchmark overview of MoCoO. (A) UMAP embeddings of all six configurations, coloured by cell type, showing qualitative clustering improvement from VAE to the Full model. (B) Radar summary of normalised metrics (ARI, NMI, ASW, DRE, DREX, LSEX) revealing that no single configuration dominates all axes. (C) Training loss convergence curves. (D) Latent space diagnostics (μ -norm, σ -mean, zero-dimension count, pairwise distance). (E) Learned ODE vector field projected onto UMAP coordinates for the Full model, visualising the continuous dynamics captured by the trajectory module. (F) Column-normalised heatmap across all evaluation metrics, providing a quantitative bird's-eye comparison.

- [5] R. Lopez, J. Regier, M. B. Cole, M. I. Jordan, and N. Yosef, "Deep generative modeling for single-cell transcriptomics," *Nature Methods*, vol. 15, no. 12, pp. 1053–1058, 2018.
- [6] C. H. Grønbech, M. F. Vording, P. N. Timshel, C. K. Sønderby, T. H. Pers, and O. Winther, "scVAE: variational auto-encoders for single-cell gene expression data," *Bioinformatics*, vol. 36, no. 16, pp. 4415–4422, 2020.
- [7] R. T. Q. Chen, Y. Rubanova, J. Bettencourt, and D. Duvenaud, "Neural ordinary differential equations," in *Advances in Neural Information Processing Systems*, vol. 31, 2018, pp. 6571–6583.
- [8] Y. Rubanova, R. T. Q. Chen, and D. Duvenaud, "Latent ordinary differential equations for irregularly-sampled time series," in *Advances in Neural Information Processing Systems*, vol. 32, 2019, pp. 5320–5330.
- [9] Z. Fu, "PanODE: Neural ODE-regularized VAE for single-cell trajectory inference," 2025, manuscript in preparation.
- [10] K. He, H. Fan, Y. Wu, S. Xie, and R. Girshick, "Momentum contrast for unsupervised visual representation learning," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2020, pp. 9729–9738.
- [11] J. Li, P. Zhou, C. Xiong, and S. Hoi, "Prototypical contrastive learning of unsupervised representations," in *Proceedings of the International Conference on Learning Representations (ICLR)*, 2021.
- [12] C. Xu, R. Lopez, E. Mehlman, J. Regier, M. I. Jordan, and N. Yosef, "Probabilistic harmonization and annotation of single-cell transcriptomics data with deep generative models," *Molecular Systems Biology*, vol. 17, no. 1, p. e9620, 2021.
- [13] A. Gayoso, Z. Steier, R. Lopez, J. Regier, K. L. Nator, A. Streets, and N. Yosef, "Joint probabilistic modeling of single-cell multi-omic data with totalVI," *Nature Methods*, vol. 18, no. 3, pp. 272–282, 2021.
- [14] V. Bergen, M. Lange, S. Peidli, F. A. Wolf, and F. J. Theis, "Generalizing RNA velocity to transient cell states through dynamical modeling," *Nature Biotechnology*, vol. 38, no. 12, pp. 1408–1414, 2020.
- [15] T. Chen, S. Kornblith, M. Norouzi, and G. Hinton, "A simple framework for contrastive learning of visual representations," in *Proceedings of the International Conference on Machine Learning (ICML)*, 2020, pp. 1597–1607.
- [16] Y. Cai, Z. Zhang *et al.*, "scAGCL: Augmentation-guided contrastive learning for single-cell RNA-seq," *Bioinformatics*, 2024, in press.
- [17] —, "scGPCL: Graph prototype contrastive learning for single-cell data," *Briefings in Bioinformatics*, 2024, in press.
- [18] T. Hao, Y. Yu *et al.*, "scDiff: Diffusion-based generative model for single-cell RNA-seq data," *Genome Biology*, 2024, in press.
- [19] Y. Li, S. Chen *et al.*, "VeloVAE: Variational autoencoder for RNA velocity," *Nature Methods*, 2024, in press.
- [20] F. Paul, Y. Arkin, A. Giladi, D. A. Jaitin, E. Kenigsberg, H. Keren-Shaul, D. Winter, D. Lara-Astiaso, M. Gury, A. Weiner, E. David, N. Cohen, F. K. B. Lauridsen, S. Haas, A. Schlitzer, A. Mildner, F. Ginhoux,

S. Jung, A. Perse, and I. Amit, "Transcriptional heterogeneity and lineage commitment in myeloid progenitors," *Cell*, vol. 163, no. 7, pp. 1663–1677, 2015.